

Fast, Effective and Interpretable Deep Learning

Frederik Hvilshøj

March 3, 2020

Table of contents


1. Introuction
2. Normalizing Flows
3. Explainability

Introuction

- Deep Learning is **leading paradigm** in Machine Learning
- Much effort put into network **design** and **efficiency**
- Models may be **biased** due to model constraints
- Limited insights – knowledge is based on **intuition**
- **Explanations** are important for sensitive tasks

This Project

- Explanations of general neural network is challenging
- Invertible functions give a higher level of insights



Normalizing Flows

Normalizing Flows

Goal: Learn a representation of a density $p(x)$

Model: Maximize log-likelihood of the data with respect to the parameters of an invertible Neural Network f .

$$\log(p_\theta(x)) = \log(p_z(f(x))) + \log\left(\left|\det \frac{\partial f(x)}{\partial x}\right|\right)$$

Constraints: invertible f and tractable Jacobian.

$$\log\left(\left|\det \frac{\partial f_k(f_{k-1}(\dots f_1(x)))}{\partial x}\right|\right) = \sum_{i=1}^k \log\left(\left|\det \frac{\partial f_i(z_{i-1})}{\partial z_{i-1}}\right|\right)$$

$$z_0 = x, \quad z_i = f_i(z_{i-1})$$

Normalizing Flow Layers

Coupling layer $c(x) = (x_1, x_2 + m(x_1))$, for $x = (x_1, x_2)$ [4]

Channel permutation Fixed channel permutations of 3D data [5]

1x1 convolutions Learnable channel interactions [9]

Periodic convolutions Circular convolutions over each channel [6]

What if Neural Networks had SVDs?

The Singular Value Decomposition

Definition

The Singular Value Decomposition of $W \in \mathbb{R}^{m \times n}$ is a factorization $W = U\Sigma V^T$, where U is an **orthogonal** $m \times m$ matrix, Σ is an $m \times n$ rectangular diagonal matrix, and V is an $n \times n$ **orthogonal** matrix.

For our purposes, assume that W is a square, i.e., $W \in \mathbb{R}^{d \times d}$. Then:

$$U^{-1} = U^T \text{ and } V^{-1} = V^T \quad O(d^2)$$

$$W^{-1} = (U\Sigma V^T)^{-1} = V\Sigma^{-1}U^T \quad O(d^2)$$

$$\det(W) = \prod_{i=1}^d \Sigma_{i,i} \quad O(d)$$

Other fast computations are largest singular value, condition number, matrix exponential, and Cayley Map.

Neural Networks

Recall a simple fully-connected layer:

$$h_\ell = \sigma(Wh_{\ell-1}) = \sigma(U\Sigma V^T h_{\ell-1})$$

- Normalizing flows
- Recurrent Neural Networks
- W-GAN

Gradients

$$W^{t+1} = W^t - \eta \frac{\partial L}{\partial W}$$

$$\Sigma^{t+1} = \Sigma^t - \eta \frac{\partial L}{\partial \Sigma}$$

$$U^{t+1} = U^t - \eta \frac{\partial L}{\partial U}$$

The Householder Matrix

Definition

For a vector $v \in \mathbb{R}^d$, the corresponding Householder matrix is defined as

$$H = I - 2 \frac{vv^T}{\|v\|_2^2} \quad (1)$$

- Note LHS of Eqn. (1) takes $O(d^2)$ but RHS takes $O(d)$.
- H stays orthogonal under gradient descent on v

$$UX = H_1(H_2(\cdots(H_dX)))$$

with $X \in \mathbb{R}^{d \times m}$, where m is the batch size.

If we represent H_i s as matrices, this takes $O(d^3m)$ time (**parallel** [10]).

We can easily improve by representing H_i by its vectors v_i instead, $O(d^2m)$ (**sequential** [10]).

Issues:

- The **parallel** algorithm is slow in theory.
- The **sequential** algorithm is not suited for GPUs.

We propose:

- Algorithm which is suited for GPUs and fast in theory.

Lemma

[3] For any m Householder matrices H_1, \dots, H_m there exists $W, Y \in \mathbb{R}^{d \times m}$ st.

$$I - 2WY^T = H_1 \cdots H_m.$$

Both W and Y can be computed by m sequential Householder multiplications in $O(dm^2)$ time.

H_1	\cdots	H_m	\cdots	H_{2m}	\cdots	H_d
-------	----------	-------	----------	----------	----------	-------

$$\underbrace{\hspace{1.5cm}}_{P_1 = I - 2W_1Y_1^T} \underbrace{\hspace{1.5cm}}_{P_2} \cdots \underbrace{\hspace{1.5cm}}_{P_{\frac{d}{m}}}$$



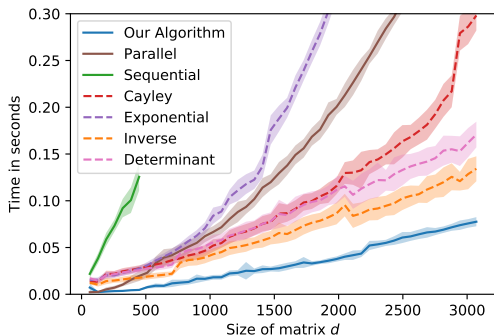
Time:

$$P_1 \cdot P_2 \cdots P_{\frac{d}{m}} X \quad O(d^3 m) \Rightarrow O(d^2 m)$$

sequential operations:


$$O(\text{compute } I - 2WY^T + \text{multiply } X \text{ with } P_i\text{'s}) = O(m + d/m)$$

Experiments



Circular 3D Convolutions

3D Circular Convolutions



Algorithm: Forward $O(hwc \log hwc)$, Determinant $O(hwc)$

Require: Image/activations $X \in \mathbb{R}^{h \times w \times c}$ and kernel $K \in \mathbb{R}^{h \times w \times c}$.

- 1: $X = \text{DFT}_{3D}(X)$
- 2: $K = \text{DFT}_{3D}(K)$
- 3: $X = X \odot K$
- 4: $X = \text{DFT}_{3D}^{-1}(X)$

Comparison with Periodic Convolutions

For $X, K \in \mathbb{R}^{h \times w \times c}$ and $W \in \mathbb{R}^{c \times c \times h \times w}$

Circular 3D convolutions

$$Z = F_{3D}^{-1}(F_{3D}(X) \odot F_{3D}(K))$$

Periodic convolutions

$$z_i = F_{2D}^{-1}\left(\sum_{j=1}^c F_{2D}(W_{i,j}) \odot F_{2D}(X_{[:, :, j]})\right),$$

Type	Params	Forward	Inverse
Periodic	$h \cdot w \cdot c^2$	$O(hwc^2)$	$O(hwc^3)$
3D-circular [†]	$h \cdot w \cdot c^2$	$O(hwc^2 \log(hwc))$	$O(hwc^2 \log(hwc))$

Improving Variational Auto-Encoders

Improving Variational Auto-Encoders

Variational lower-bound:

$$\log p(x) \geq \mathbb{E}_{z \sim q(z|x)} [\log p(x|z)] + D_{KL}(q(z|x) || p(z))$$

where $q(z|x) = \mathcal{N}(\mu(x), \text{diag}(\sigma(x)))$.

With formalizing flow f , the equation becomes

$$\log p(x) \geq \mathbb{E}_{z \sim p(z|x)} \left[\log p(x|f(z)) + \log \left| \text{Det} \frac{\partial f(z)}{\partial z} \right| \right] + D_{KL}(q(z|x) || p(f(z)))$$

[11] let $f(z) = H_T H_{T-1} \cdots H_1 z$ with each H_i being parametrized by the encoder.

Improving Variational Auto-Encoders

Method	$\leq \ln p(\mathbf{x})$
VAE	-93.89 ± 0.09
VAE+HF($T=1$)	-87.77 ± 0.05
VAE+HF($T=10$)	-87.68 ± 0.06
VAE+NF ($T=10$) [17]	-87.5
VAE+NF ($T=80$) [17]	-85.1
VAE+NICE ($T=10$) [7]	-88.6
VAE+NICE ($T=80$) [7]	-87.2
VAE+HVI ($T=1$) [20]	-91.70
VAE+HVI ($T=8$) [20]	-88.30

Figure 1: Comparison of the lower bound of marginal log-likelihood measured in nats of the digits in the MNIST test set. For the first three methods the experiment was repeated 3 times. Direct copy from [11].

Explainability

Here we focus on **explainability**, characterized by

An active characteristic of a model, denoting any action or procedure taken by a model with the intent of clarifying or detailing its internal functions [1].

We consider **attribution** techniques and **counterfactual** explanations.

Attribution Techniques

- Identifies contribution of each input feature wrt. the output
- “Gradient-like” computations
- Based on different propagation rules




Figure 2: Direct copy from [8].

Counterfactual Explanations

Counterfactual explanations tries to answer the question

*How can I make a **minimal** and **realistic** change to an input of the model such that the predicted outcome changes?*

Common strategy

1. Start from the original input
2. Until prediction change
 - 2.1 Make small change to the input
 - 2.2 Query the classifier to gain information

Differences

- How to do update
- Information gained from classifier


Improving Explanations with Probabilistic Saliency Estimation

Probabilistic Saliency Estimation [2]


Input ($X \in \mathbb{R}^{h \times w \times c}$)



Super-pixels




$P \in \mathbb{R}^{h \times w \times c}$



Let $A \in \mathbb{R}^{h \times w \times c}$ be an attribution map, then we propose a refined attribution

$$\hat{A} = A \odot P$$

Experiment



- Fill-methods
- Evidence against classes
- Larger tests

Attention Mechanisms in Gradient Explanations

Guiding backward computations

A simple dense layer $f_\ell : L' \rightarrow L$ of a Neural Network with backpropagation function $g_\ell : L \rightarrow L'$ of explanation a_ℓ is

$$f_\ell(h_{\ell-1}) = \sigma(h_{\ell-1}^T W_\ell) \qquad g_\ell(a_\ell) = a_\ell^T \frac{\partial \sigma(h_{\ell-1} W_\ell)}{\partial h_{\ell-1} W_\ell} W_\ell^T$$

with $W \in \mathbb{R}^{L' \times L}$. We experimented with guiding the backpropagation

$$\begin{aligned} \hat{a}_{\ell-1} &= g_\ell(a_\ell) \odot \alpha_\ell(h_{\ell-1}, a_\ell) \\ \alpha_\ell(h_{\ell-1}, a_\ell) &= \tanh([h_{\ell-1} || a_\ell] L_\ell R_\ell) \end{aligned}$$

with $L_\ell \in \mathbb{R}^{(L'+L) \times d}$ and $R_\ell \in \mathbb{R}^{d \times L'}$, for $d < \frac{L \cdot L'}{L+L'}$

Extensions:

- Change the input to the α_ℓ s
- Scale the output to be in range $[-c, c]$

- Need to define a loss
- DeconvNet and Guided Backprop are class insensitive

$$L(X, \hat{X}) = \|\hat{X}\|_1 + \lambda \mathcal{L}(y, \hat{y}),$$

Extensions:

$$L(X, \hat{X}) = \mathcal{L}(y, \hat{y}) + \lambda_1 \|\hat{X}\|_1 + \lambda_2 \|X - \hat{X}\|_2 + \lambda_3 \text{vec}(\hat{X})^T L \text{vec}(\hat{X})$$

$$f^T L f = \frac{1}{2} \sum_{i=1}^d \sum_{j=1}^d a_{i,j} (f_i - f_j)^2$$


Challenges

- Currently only potentially works with DeconvNet, GBP, and PatternNet
- Loss for attribution methods?
- Are these explanations from the in-sample-distribution?

$$\text{AOPC} = \frac{1}{L+1} \left\langle \sum_{k=0}^L f(x_{\text{MoRF}}^{(0)}) - f(x_{\text{MoRF}}^{(k)}) \right\rangle_{p(x)}$$

Generating Counterfactual Explanations


Construction



$$f(x) = \sigma(\beta^T NN(x))$$

$$f^\dagger(y) = NN^{-1} \left(f(x) - \frac{\langle f(x), \beta \rangle \beta}{\|\beta\|} \right)$$

Properties and Extensions



Properties

- First “one-evaluation-solution” (Fast)
- Built to produce counterfactual explanations

Extensions

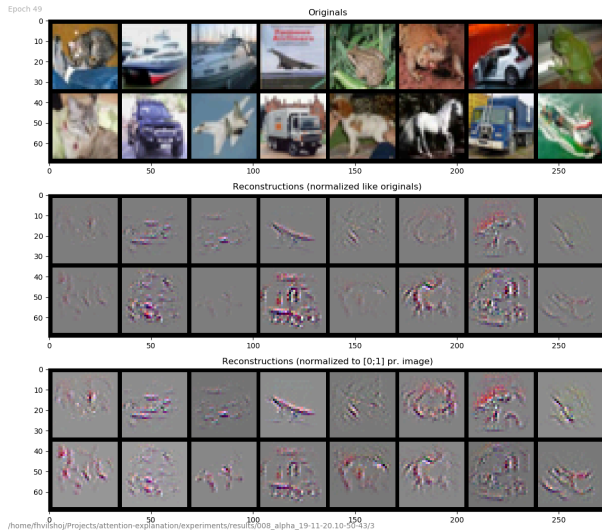
- Training NN and β jointly
- Evaluating Counterfactual explanations
- Bound difference in X from change in Z space
 - Lipschitz constraint on NN

$$Score(X, \hat{X}) = \|\Delta_X\|_1 + \text{vec}(\Delta_X) L \text{vec}(\Delta_X) + \mathbf{1}_{f(\hat{X}) \neq \bar{f}(\hat{X})}$$

- [1] Alejandro Barredo Arrieta, Natalia Díaz-Rodríguez, Javier Del Ser, Adrien Bennetot, Siham Tabik, Alberto Barbado, Salvador García, Sergio Gil-López, Daniel Molina, Richard Benjamins, Raja Chatila, and Francisco Herrera. Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities and Challenges toward Responsible AI. *arXiv preprint arXiv:1910.10045*, 2019.
- [2] Caglar Aytekin, Alexandros Iosifidis, and Moncef Gabbouj. Probabilistic saliency estimation. *Pattern Recognition*, 74:359–372, 2018.
- [3] Christian Bischof and Charles Van Loan. The WY Representation for Products of Householder Matrices. *SIAM Journal on Scientific and Statistical Computing*, 1987.
- [4] Laurent Dinh, David Krueger, and Yoshua Bengio. NICE: Non-Linear Independent Components Estimation. In *ICLR (Workshop)*, 2015.
- [5] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using real NVP. In *ICLR*, 2019.
- [6] Emiel Hoogeboom, Rianne van den Berg, and Max Welling. Emerging Convolutions for Generative Normalizing Flows. In *ICML*, 2019.

- [7] Jörn Henrik Jacobsen, Jens Behrmann, Richard Zemel, and Matthias Bethge. Excessive invariance causes adversarial vulnerability. In *ICLR*, 2019.
- [8] Pieter Jan Kindermans, Kristof T. Schütt, Maximilian Alber, Klaus Robert Müller, Dumitru Erhan, Been Kim, and Sven Dähne. Learning how to explain neural networks: Patternnet and Patternattribution. In *ICLR*, 2018.
- [9] Diederik P. Kingma and Prafulla Dhariwal. Glow: Generative Flow with Invertible 1x1 Convolutions. In *NeurIPS*, 2018.
- [10] Zakaria Mhammedi, Andrew Hellicar, Ashfaqur Rahman, and James Bailey. Efficient Orthogonal Parametrisation of Recurrent Neural Networks Using Householder Reflections. In *ICML*, 2017.
- [11] Jakub M. Tomczak and Max Welling. Improving Variational Auto-Encoders using Householder Flow. *arXiv preprint arXiv:1611.09630*, 2016.

Attention example



Attention example




Figure 3: Direct copy from [7].